

LÁZARO GÓMEZ VITOLO

Argentina

AI & Robotics Developer | Python • Machine Learning • Computer Vision • Generative AI

Email: lazarogvitolo@gmail.com • GitHub: [Lazaro549](#) • Portfolio: [lazaro549.github.io/Portfolio/](#) •

LinkedIn: [lazaro-gomez-vitolo](#)

PROFESSIONAL PROFILE

AI-focused software developer building practical, production-ready systems across Generative AI, LLM evaluation, RAG architectures, computer vision pipelines, and robotics control interfaces. Strong commitment to reproducible experiments, measurable benchmarks, modular architecture, and robust engineering workflows. Public portfolio demonstrates technical depth across AWS Bedrock, Anthropic Claude API, Model Context Protocol (MCP), structured evaluation pipelines, and browser-based real-time perception.

CORE TECHNICAL SKILLS

AI & GENERATIVE MODELS LLM evaluation, LLM-as-a-judge, RAG, embeddings, benchmarking, prompt engineering, AI agents, ML workflows	ROBOTICS & COMPUTER VISION Object detection, real-time perception, tracking algorithms, telemetry concepts, CV pipelines, browser inference
SOFTWARE ENGINEERING Python, TypeScript, JavaScript, FastAPI, React, Vite, Streamlit, RESTful APIs, SQLite, Docker, GitHub Actions, pytest	CLOUD & AI INFRASTRUCTURE Amazon Bedrock, AWS Lambda, AWS SAM, Anthropic Claude API, MCP (Model Context Protocol), OpenAI API workflows

SELECTED ENGINEERING PROJECTS

AI Evaluation Platform Stack: Python • FastAPI • React • TypeScript • SQLite • Docker • GitHub Actions Production-oriented evaluation platform for LLM applications. Implements deterministic and LLM-as-a-judge metrics, batch evaluation, provider abstraction (Mock/OpenAI/Bedrock), run comparison tools, automated reporting, REST API, and dashboard with API-key-free CI integration.
AWS AI Generative / RAG Benchmark Stack: AWS Bedrock • FAISS • Python • RAG • AWS Lambda • SAM End-to-end AWS generative AI repository featuring a chatbot, RAG pipeline, serverless API, and evaluation framework. Controlled retrieval experiment evaluated chunk size, overlap, and top-k across 6 documents and 18 questions; balanced configuration boosted recall from 0.250 to 0.667 while lowering retrieval latency.
Claude Code / MCP Engineering Workflows Stack: Claude Code • TypeScript • MCP • GitHub Actions Architected reproducible examples for Claude Code skills, subagents, custom hooks, and CI/CD pipelines. Developed an advanced MCP server covering sampling, elicitation, roots, pagination, structured output, Streamable HTTP, session resumability, and OAuth 2.1 + PKCE authentication.
Robot Vision System Stack: TensorFlow.js • COCO-SSD • JavaScript • Vitest Browser-based real-time perception dashboard using webcam and video input, object detection, confidence filtering, centroid tracking, and FPS/latency metrics. Explicitly simulates robot battery, odometry, and telemetry data.
AI Engineering Lab Stack: TypeScript • Vite • React Component Architecture • GitHub Pages Technical hub linking build, evaluate, benchmark, improve, and ship workflows. Enforces evidence integrity by strictly separating measured, demo, and unverified results with tests that prevent fabricated metrics.

CERTIFICATIONS & TRAINING

- **Anthropic Academy:** Claude Code in Action
- **Anthropic Academy:** Building with the Claude API (Certificate Project)
- **Anthropic Academy:** Model Context Protocol: Advanced Topics
- **OpenAI:** Agents & Workflows